


Disk Based Data Structures

- So far search trees were limited to main memory structures
 - Assumption: the dataset organized in a search tree fits in main memory (including the tree overhead)
- Counter-example: transaction data of a bank > 1 GB per day
 - use secondary storage media (punch cards, hard disks, magnetic tapes, etc.)
- Consequence: make a search tree structure secondary-storage-enabled



Hard Disks

- ❑ Large amounts of storage, but slow access!
- ❑ Identifying a page takes a long time (seek time plus rotational delay – 5-10ms), reading it is fast
 - ❑ It pays off to read or write data in **pages** (or blocks) of 2-16 Kb in size.



Algorithm analysis

- The running time of disk-based algorithms is measured in terms of
 - computing time (CPU)
 - number of disk accesses
 - sequential reads
 - random reads
- Regular main-memory algorithms that work one data element at a time can not be “ported” to secondary storage in a straight-forward way

Principles

- Pointers in data structures are no longer addresses in main memory but locations in files
- If x is a pointer to an object
 - if x is in main memory $key[x]$ refers to it
 - otherwise $DiskRead(x)$ reads the object from disk into main memory ($DiskWrite(x)$ – writes it back to disk)

Principles (2)

□ A typical working pattern


```
01 ...  
02 x ← a pointer to some object  
03 DiskRead(x)  
04 operations that access and/or modify x  
05 DiskWrite(x) //omitted if nothing changed  
06 other operations, only access no modify  
07 ...
```

□ Operations:

- DiskRead(x:pointer_to_a_node)
- DiskWrite(x:pointer_to_a_node)
- AllocateNode():pointer_to_a_node

Binary-trees vs. B-trees

- Size of B-tree nodes is determined by the page size. One page – one node.
- A B-tree of height 2 may contain > 1 billion keys!
- Heights of Binary-tree and B-tree are logarithmic
 - B-tree: logarithm of base, e.g., 1000
 - Binary-tree: logarithm of base 2



1 node
1000 keys

1001 nodes,
1,001,000 keys

1,002,001 nodes,
1,002,001,000 keys

B-tree Definitions

- Node x has fields
 - $n[x]$: the number of keys of that the node
 - $\text{key}_1[x] \leq \dots \leq \text{key}_{n[x]}[x]$: the keys in ascending order
 - $\text{leaf}[x]$: true if leaf node, false if internal node
 - if internal node, then $c_1[x], \dots, c_{n[x]+1}[x]$: pointers to children
- Keys separate the ranges of keys in the subtrees. If k_i is an arbitrary key in the subtree $c_i[x]$ then $k_i \leq \text{key}_i[x] \leq k_{i+1}$


B-tree Definitions (2)

- Every leaf has the same depth
- In a B-tree of a **degree** t all nodes except the root node have between t and $2t$ children (i.e., between $t-1$ and $2t-1$ keys).
- The root node has between 0 and $2t$ children (i.e., between 0 and $2t-1$ keys)

Height of a B-tree

- B-tree T of height h , containing $n \geq 1$ keys and minimum degree $t \geq 2$, the following restriction on the height holds:

$$h \leq \log_t \frac{n+1}{2}$$



$$n \geq 1 + (t-1) \sum_{i=1}^h 2t^{i-1} = 2t^h - 1$$

B-tree Operations

- An implementation needs to support the following B-tree operations
 - **Searching** (simple)
 - **Creating** an empty tree (trivial)
 - **Insertion** (complex)
 - **Deletion** (complex)

Searching

- Straightforward generalization of a binary tree search

```
BTreeSearch(x,k)
01 i ← 1
02 while i ≤ n[x] and k > keyi[x]
03     i ← i+1
04 if i ≤ n[x] and k = keyi[x] then
05     return(x,i)
06 if leaf[x] then
07     return NIL
08 else DiskRead(ci[x])
09     return BTtreeSearch(ci[x],k)
```

Creating an Empty Tree

- Empty B-tree = create a root & write it to disk!

BTreeCreate(T)


```
01 x ← AllocateNode();  
02 leaf[x] ← TRUE;  
03 n[x] ← 0;  
04 DiskWrite(x);  
05 root[T] ← x
```

Splitting Nodes

- Nodes fill up and reach their maximum capacity $2t - 1$
- Before we can insert a new key, we have to “make room,” i.e., split nodes

Splitting Nodes (2)

- Result: one key of x moves up to parent + 2 nodes with $t-1$ keys



Splitting Nodes (2)


```
BTreeSplitChild(x,i,y)
  z ← AllocateNode()
  leaf[z] ← leaf[y]
  n[z] ← t-1
  for j ← 1 to t-1
    keyj[z] ← keyj+t[y]
  if not leaf[y] then
    for j ← 1 to t
      cj[z] ← cj+t[y]
  n[y] ← t-1
  for j ← n[x]+1 downto i+1
    cj+1[x] ← cj[x]
  ci+1[x] ← z
  for j ← n[x] downto i
    keyj+1[x] ← keyj[x]
  keyi[x] ← keyt[y]
  n[x] ← n[x]+1
  DiskWrite(y)
  DiskWrite(z)
  DiskWrite(x)
```

x: parent node

y: node to be split and child of x

i: index in x

z: new node



Split: Running Time

- A local operation that does not traverse the tree
- $\Theta(t)$ CPU-time, since two loops run t times
- 3 I/Os

Inserting Keys

- Done recursively, by starting from the root and recursively traversing down the tree to the leaf level
- Before descending to a lower level in the tree, make sure that the node contains $< 2t - 1$ keys:
 - so that if we split a node in a lower level we will have space to include a new key


Inserting Keys (2)

- Special case: root is full (BtreeInsert)

```
BTreeInsert(T)
  r ← root[T]
  if n[r] = 2t - 1 then
    s ← AllocateNode()
    root[T] ← s
    leaf[s] ← FALSE
    n[s] ← 0
    c1[s] ← r
    BTreeSplitChild(s, 1, r)
    BTreeInsertNonFull(s, k)
  else BTreeInsertNonFull(r, k)
```

Splitting the Root

- Splitting the root requires the creation of a new root



- The tree grows at the top instead of the bottom

Inserting Keys

- BtreeNonFull tries to insert a key k into a node x , which is **assumed to be non-full** when the procedure is called
- BTreeInsert and the recursion in BTreeInsertNonFull guarantee that this assumption is true!

Inserting Keys: Pseudo Code

BTreeInsertNonFull(x, k)

```
01  $i \leftarrow n[x]$ 
02 if leaf[ $x$ ] then
03     while  $i \geq 1$  and  $k < \text{key}_i[x]$ 
04          $\text{key}_{i+1}[x] \leftarrow \text{key}_i[x]$ 
05          $i \leftarrow i - 1$ 
06      $\text{key}_{i+1}[x] \leftarrow k$ 
07      $n[x] \leftarrow n[x] + 1$ 
08     DiskWrite( $x$ )
```


leaf insertion

```
09 else while  $i \geq 1$  and  $k < \text{key}_i[x]$ 
10      $i \leftarrow i - 1$ 
11      $i \leftarrow i + 1$ 
12     DiskRead  $c_i[x]$ 
13     if  $n[c_i[x]] = 2t - 1$  then
14         BTreeSplitChild( $x, i, c_i[x]$ )
15         if  $k > \text{key}_i[x]$  then
16              $i \leftarrow i + 1$ 
17     BTreeInsertNonFull( $c_i[x], k$ )
```


internal node:
traversing tree

Insertion: Example


initial tree ($t = 3$)



B inserted




Q inserted




Insertion: Example (2)

L inserted



F inserted



Insertion: Running Time


- Disk I/O: $O(h)$, since only $O(1)$ disk accesses are performed during recursive calls of `BTreeInsertNonFull`
- CPU: $O(th) = O(t \log_t n)$
- At any given time there are $O(1)$ number of disk pages in main memory

Deleting Keys


- Done recursively, by starting from the root and recursively traversing down the tree to the leaf level
- Before descending to a lower level in the tree, make sure that the node contains $\geq t$ keys (cf. insertion $< 2t - 1$ keys)
- BtreeDelete distinguishes three different stages/scenarios for deletion
 - Case 1: key k found in leaf node
 - Case 2: key k found in internal node
 - Case 3: key k suspected in lower level node

Deleting Keys (2)

initial tree



F deleted:
case 1





- Case 1: If the key k is in node x , and x is a leaf, delete k from x

Deleting Keys (3)

- Case 2: If the key k is in node x , and x is not a leaf, delete k from x
 - a) If the child y that precedes k in node x has at least t keys, then find the predecessor k' of k in the sub-tree rooted at y . Recursively delete k' , and replace k with k' in x .
 - b) Symmetrically for successor node z

M deleted:
case 2a








Deleting Keys (4)

- If both y and z have only $t-1$ keys, **merge** k with the contents of z into y , so that x loses both k and the pointers to z , and y now contains $2t-1$ keys. Free z and recursively delete k from y .

G deleted:
case 2c





$$y = y+k + z - k$$


Deleting Keys - Distribution

- Descending down the tree: if k not found in current node x , find the sub-tree $c_i[x]$ that has to contain k .
- If $c_i[x]$ has only $t - 1$ keys take action to ensure that we descent to a node of size at least t .
- We can encounter two cases.
 - If $c_i[x]$ has only $t-1$ keys, but a sibling with at least t keys, give $c_i[x]$ an extra key by moving a key from x to $c_i[x]$, moving a key from $c_i[x]$'s immediate left and right sibling up into x , and moving the appropriate child from the sibling into $c_i[x]$ - ***distribution***

Deleting Keys – Distribution(2)




B deleted:




Deleting Keys - Merging

- If $c_i[x]$ and both of $c_i[x]$'s siblings have $t - 1$ keys, **merge** c_i with one sibling, which involves moving a key from x down into the new merged node to become the median key for that node




Deleting Keys – Merging (2)

delete **D**



D deleted:



tree shrinks in height

Deletion: Running Time

- Most of the keys are in the leaf, thus deletion most often occurs there!
- In this case deletion happens in one downward pass to the leaf level of the tree
- Deletion from an internal node might require “backing up” (case 2)
- Disk I/O: $O(h)$, since only $O(1)$ disk operations are produced during recursive calls
- CPU: $O(th) = O(t \log_t n)$

Two-pass Operations

- Simpler, practical versions of algorithms use two passes (down and up the tree):
 - *Down* – Find the node where deletion or insertion should occur
 - *Up* – If needed, split, merge, or distribute; propagate splits, merges, or distributes up the tree
- To avoid reading the same nodes twice, use a buffer of nodes